

# Python for Data Science

Kylie A. Bemis

Northeastern University  
Khoury College of Computer Sciences



Northeastern University

# Goals for today

- Data frames with **pandas**
- Visualization with **matplotlib**
- Modeling with **statsmodels** and **sklearn**
- Review (Q&A)

# SERIES AND DATA FRAMES

# Working with data in **pandas**

- Structures for representing *tabular data*
  - ◆ **Series** represent *homogenous* data attributes (vectors)
  - ◆ **DataFrames** represent tabular data (rows and columns)
- Methods for manipulating data tables
  - ◆ Subsetting, grouping, and summarization
  - ◆ Joining and organizing multiple tables
- Built on top of **NumPy**

# A **Series** represents a data vector

- *Homogenous* vector consisting of:
  - ◆ **Values** in the data series
  - ◆ **Index** of labels for each data element
  - ◆ Data type (**dtype**) of the values
- Built from a `numpy.ndarray`

# A **DataFrame** represents tabular data

- Data table made of rows and columns
  - ◆ Each *column* is represented by a **Series**
  - ◆ Columns may have *different* data types (**dtype**)
  - ◆ Columns must *share* the same **index** labels
  - ◆ Each *row* is identified by an **index** label
- Similar to a **dict** of data columns

# Creating a DataFrame from a dict

```
In : import pandas as pd
```

```
In : A = pd.DataFrame(  
    {'x': [100, 200, 300, 400],  
     'y': [1.11, 2.22, 3.33, 4.44],  
     'z': ['foo', 'bar', 'baz', 'qux']},  
    index=['a', 'b', 'c', 'd'])
```

```
In : A
```

```
Out:
```

	x	y	z
a	100	1.11	foo
b	200	2.22	bar
c	300	3.33	baz
d	400	4.44	qux

# Advanced pandas

- Grouping and aggregation
  - ◆ Summary statistics over rows or columns
  - ◆ Aggregation over categories in a column
- Merging and other relational operations
  - ◆ Join two or more tables based on common rows
  - ◆ Similar to SQL operations on a database
- Input to other data science libraries



# VISUALIZATION

# Visualization with **matplotlib**

- Provides MATLAB-like plotting interface
  - ◆ Flexible object-oriented plotting
  - ◆ Figures organized into hierarchical structure
  - ◆ Support for different GUI backends
- Export to many formats (PDF, PNG, etc.)

# Plotting a basic scatter plot

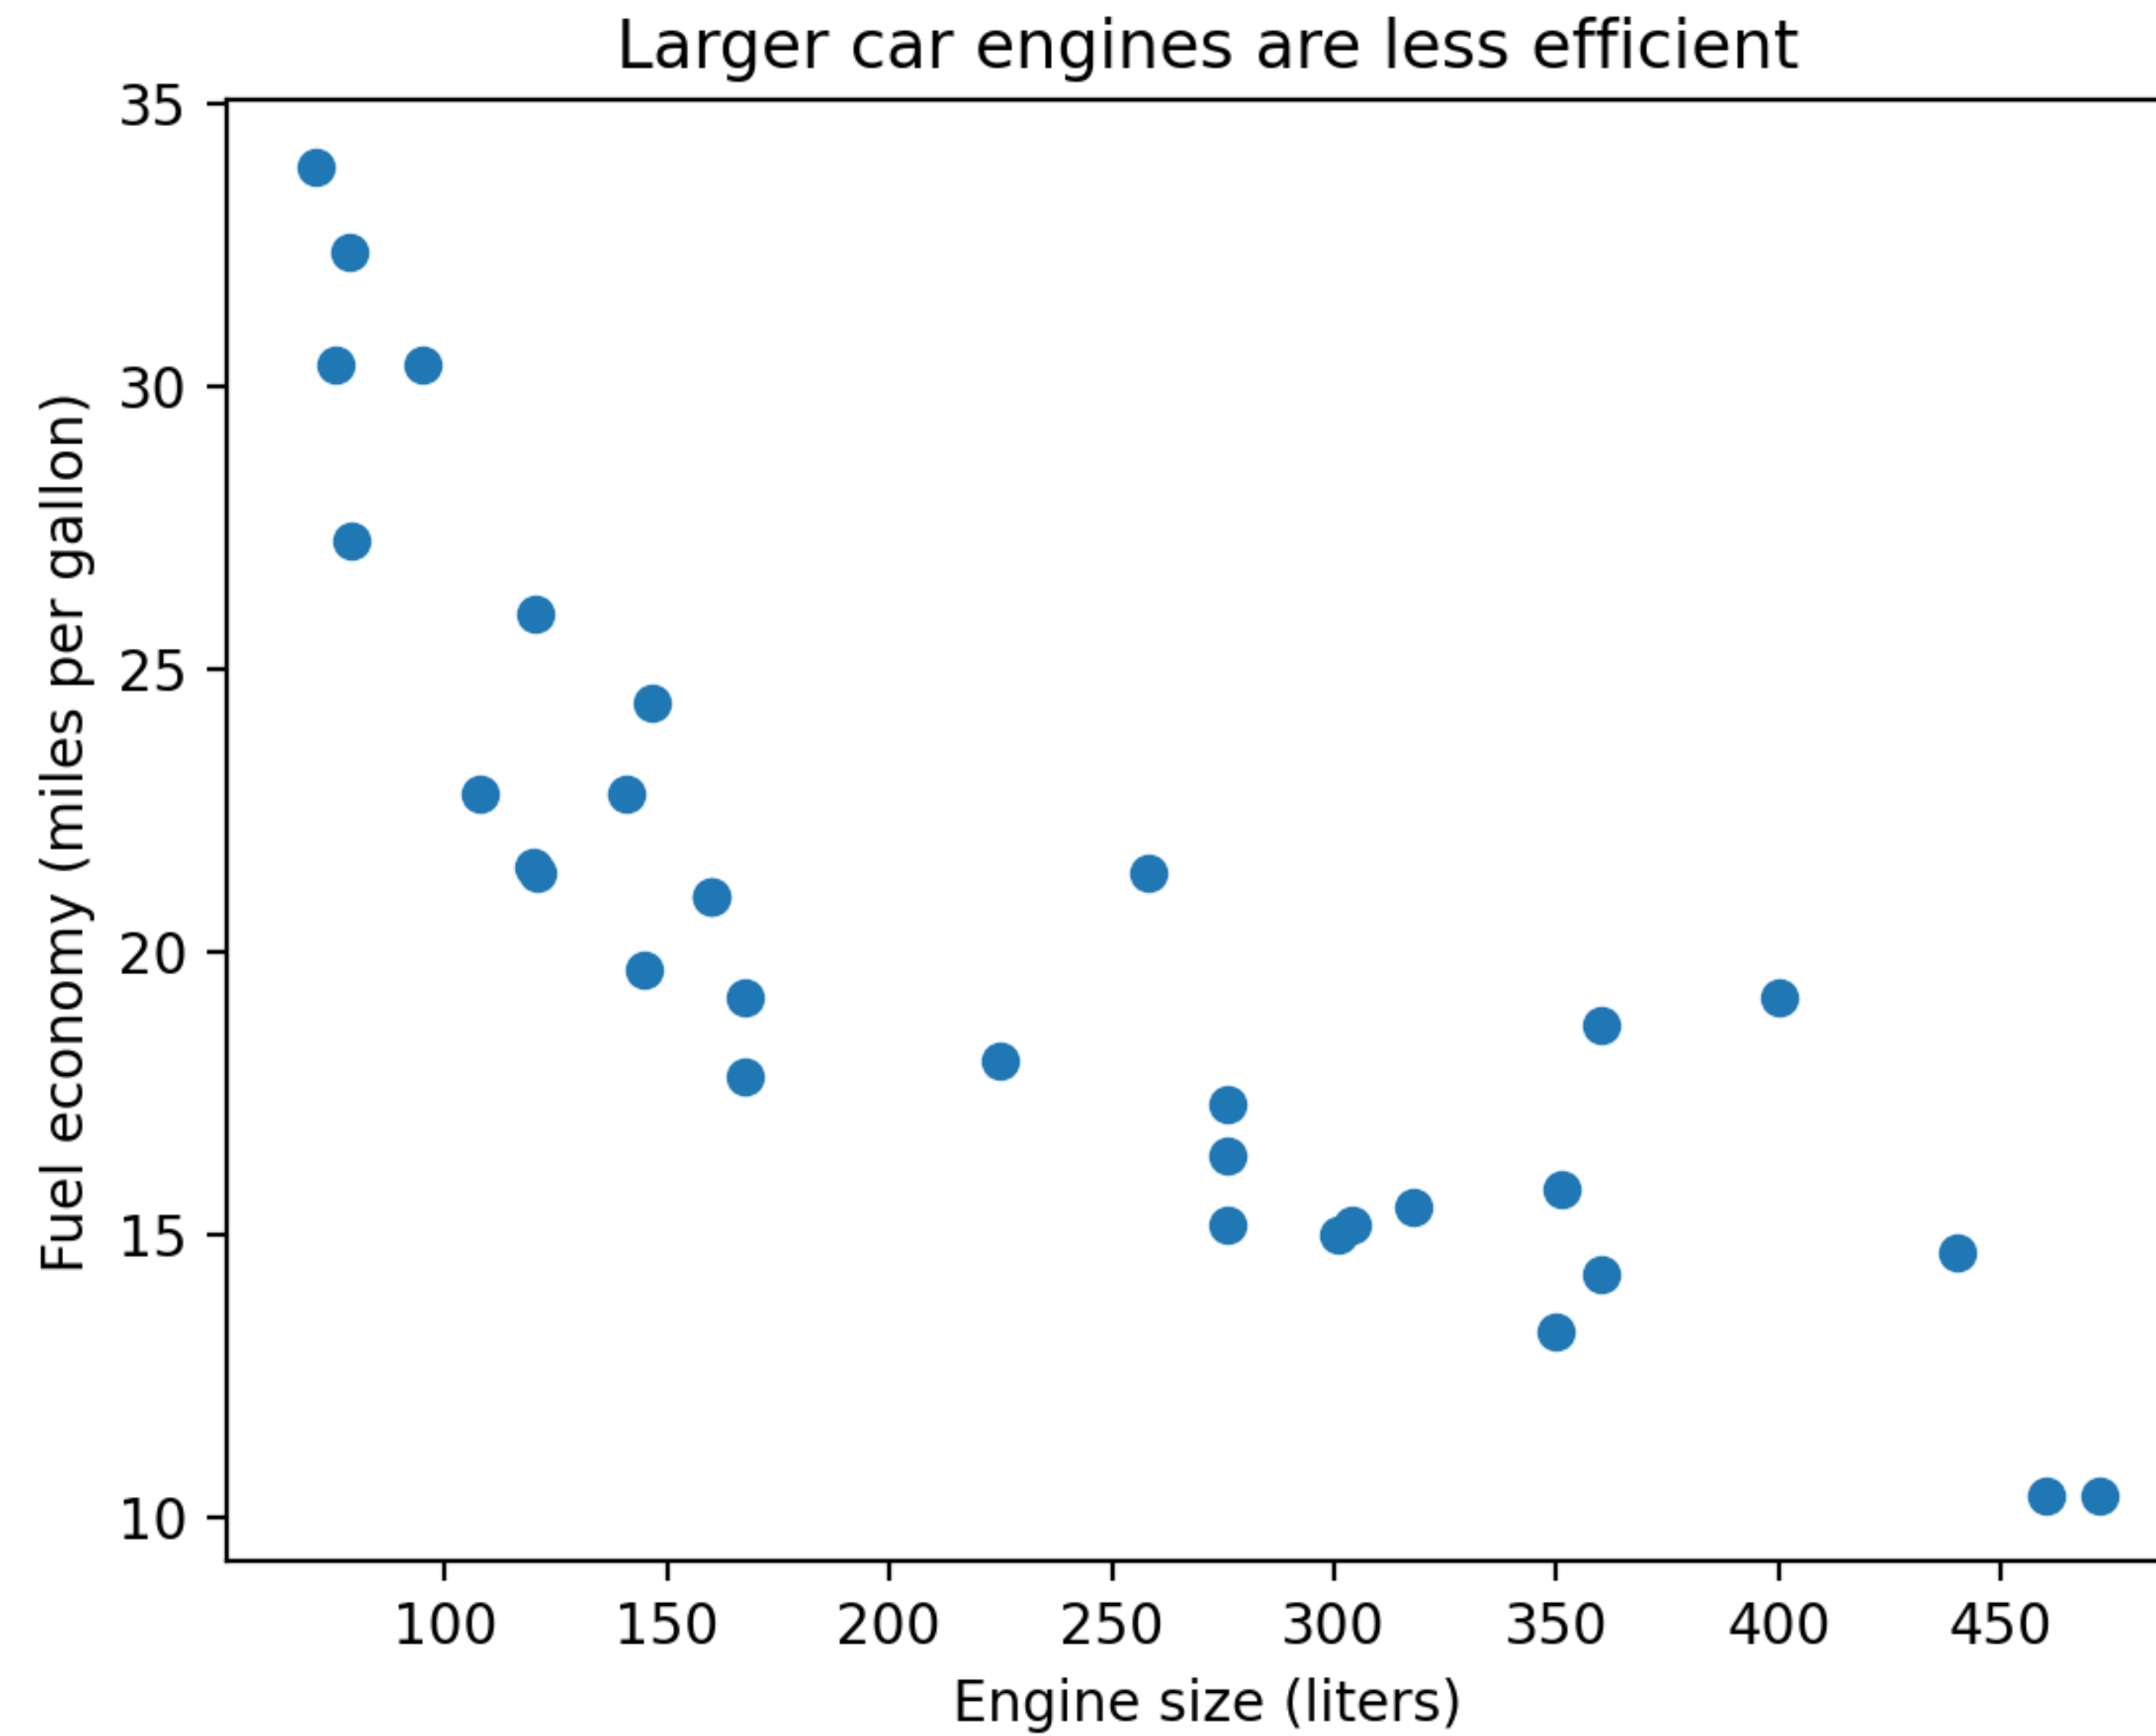
```
import matplotlib.pyplot as plt

cars = pd.read_csv("mtcars.csv", index_col=0)

xs = cars['disp']
ys = cars['mpg']

plt.scatter(xs, ys)
plt.title("Larger car engines are less efficient")
plt.xlabel("Engine size (liters)")
plt.ylabel("Fuel economy (miles per gallon)")
plt.show()
```


# Plotting a basic scatter plot (2)



# Hierarchy of a figure

- A **Figure** is *container* for other graphics
  - ◆ Outermost layer of a matplotlib graphic
  - ◆ Contains one or more Axes objects
- **Axes** objects contain *subplots*
  - ◆ Contains an individual plot or graphic
  - ◆ Does not refer to a traditional plot “axis”

# Figures and Axes



<https://towardsdatascience.com/what-are-the-plt-and-ax-in-matplotlib-exactly-d2cf4bf164a9>

# Plotting a figure with subplots

```
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2)

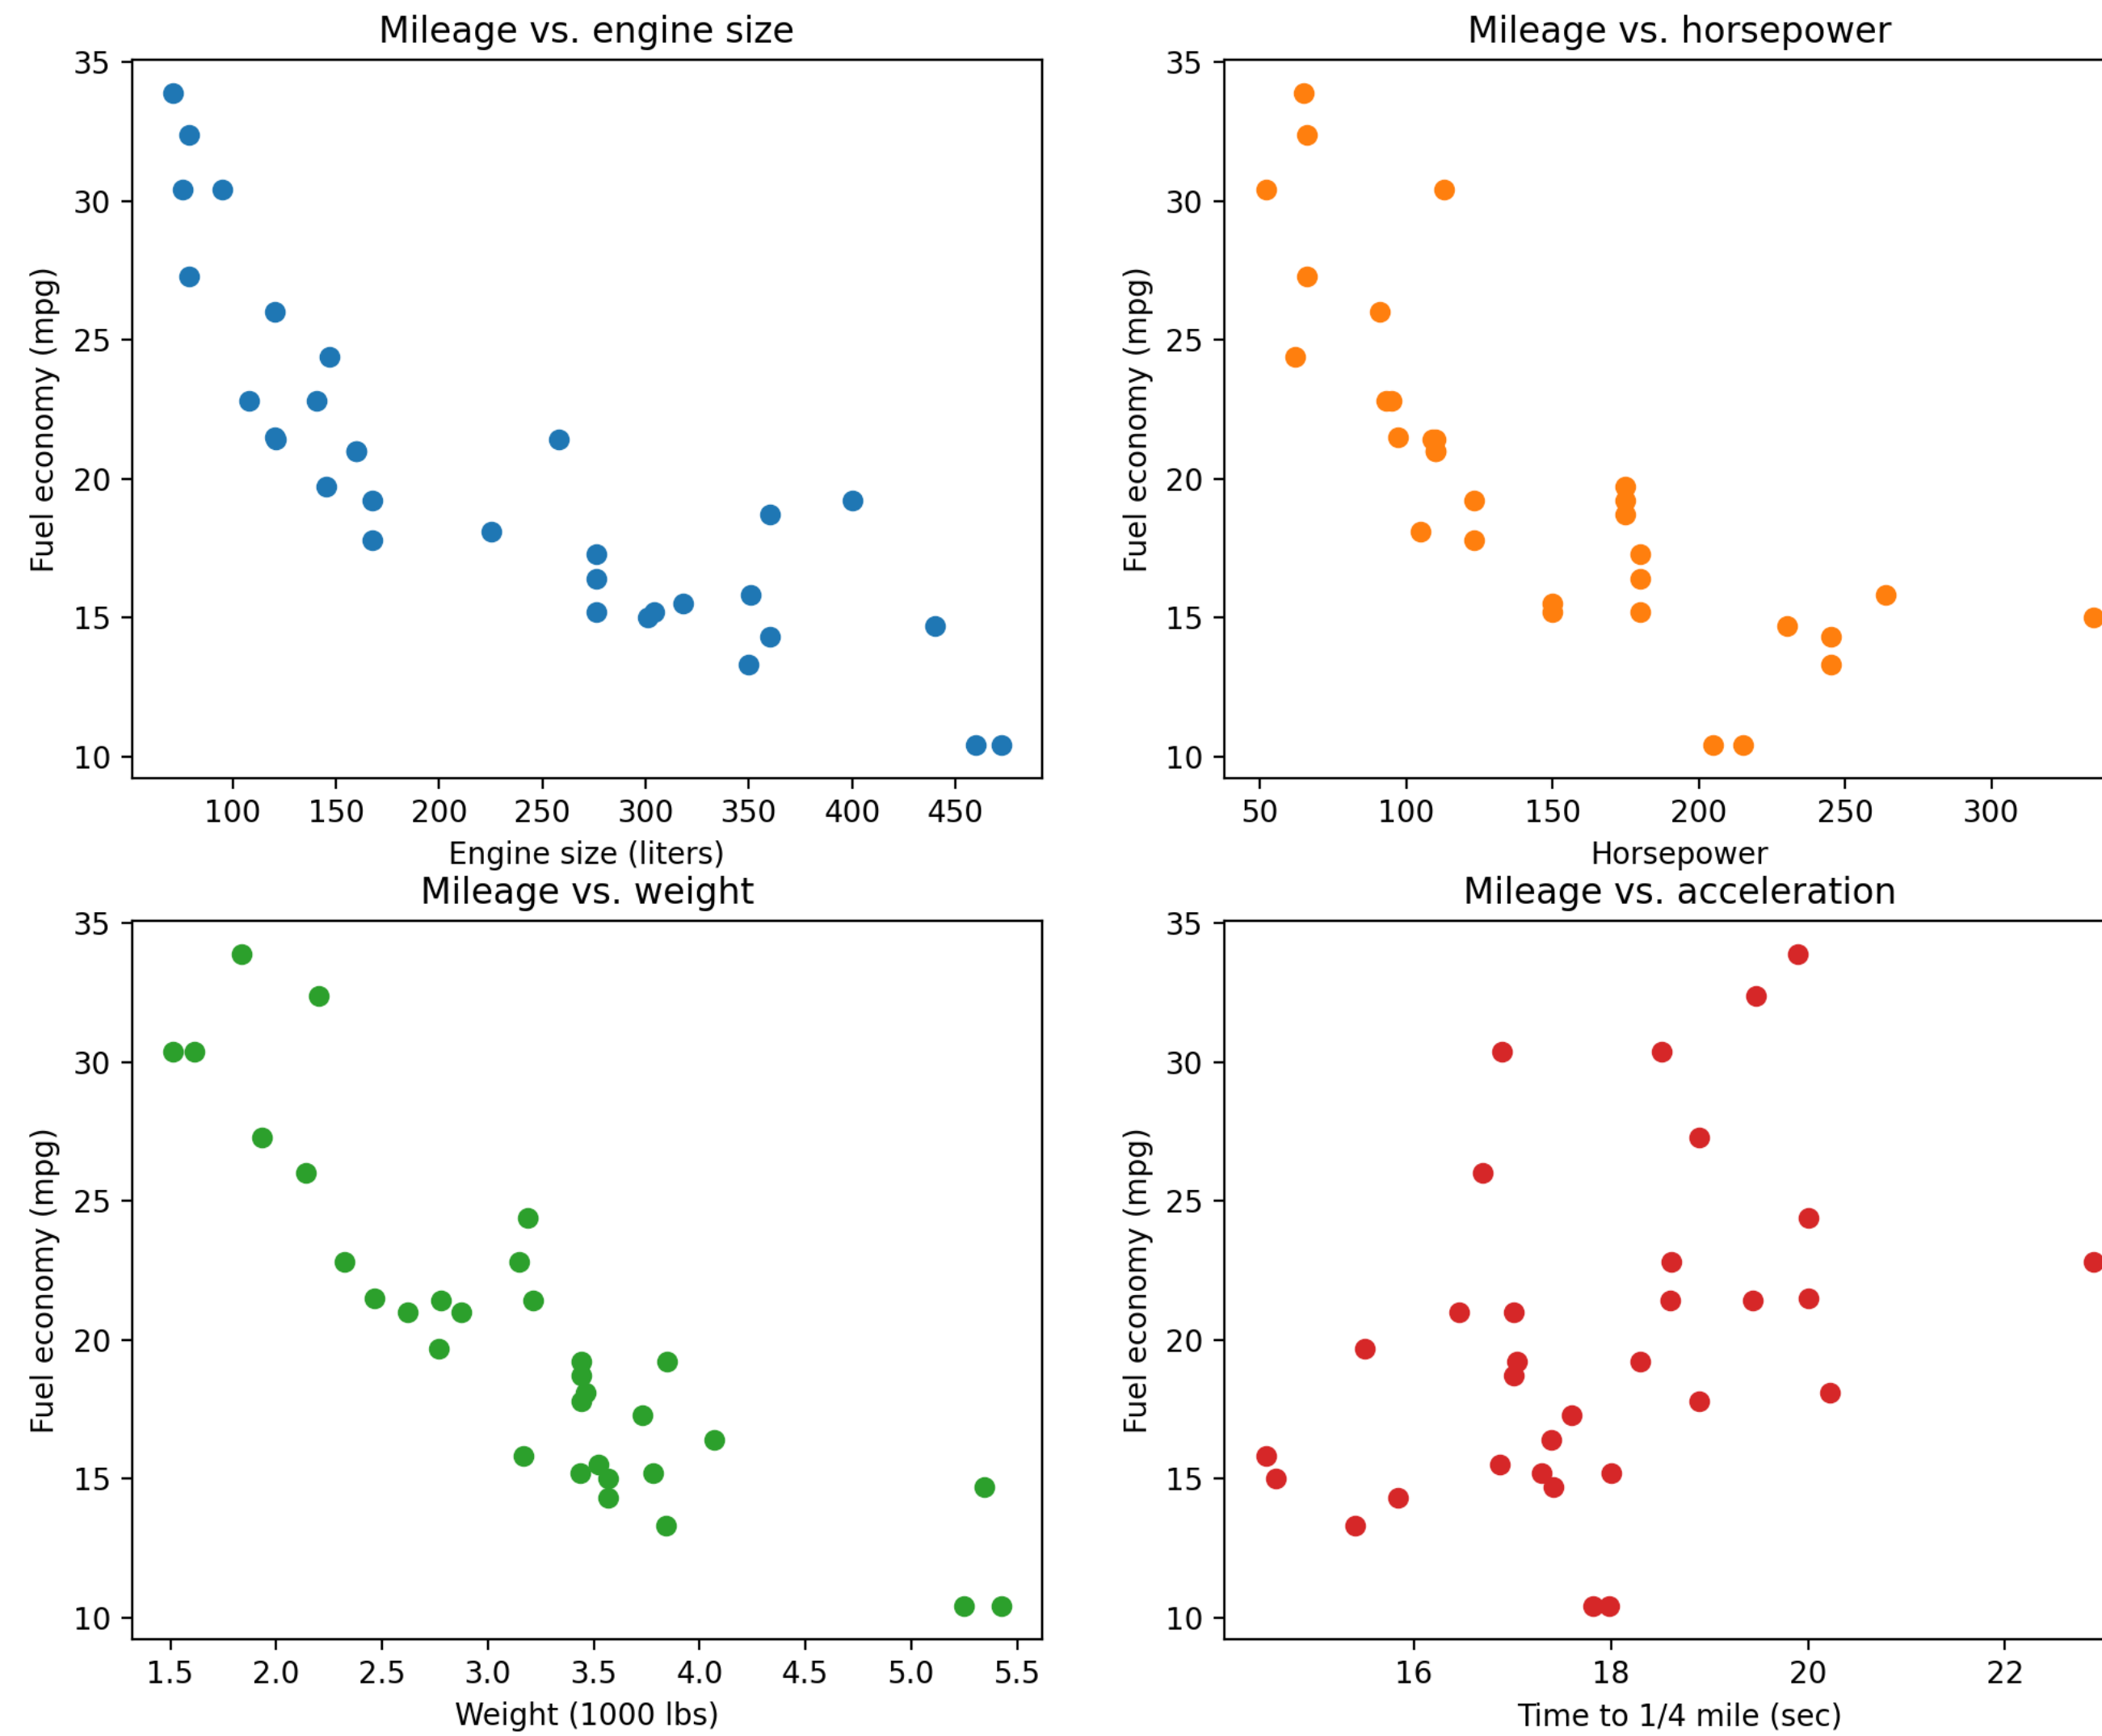
ax1.scatter('disp', 'mpg', data=cars, color='tab:blue')
ax1.set_title("Mileage vs. engine size")
ax1.set_xlabel("Engine size (liters)")
ax1.set_ylabel("Fuel economy (mpg)")

...

ax4.scatter('qsec', 'mpg', data=cars, color='tab:red')
ax4.set_title("Mileage vs. acceleration")
ax4.set_xlabel("Time to 1/4 mile (sec)")
ax4.set_ylabel("Fuel economy (mpg)")

fig.show()
```


# Plotting a figure with subplots (2)





# Figure components

- Title
- Labels
- Axes
- Ticks
- Legend
- Grid
- Spines
- Markers



# More on visualization

- **Always label your figures!!!**
  - ◆ Label the axes (include units)
  - ◆ Provide useful and informative titles
  - ◆ Include a legend whenever necessary
- Consider higher-level libraries
  - ◆ **Seaborn** provides higher-level plotting interface
  - ◆ Most built on top of **matplotlib**

# MODELING

# Modeling in Python

- **statsmodels** provides statistical models
  - ◆ More focus on statistical inference
  - ◆ More similar to other statistical software (R, Stata, etc.)
- **scikit-learn** performs machine learning
  - ◆ More focus on predictive performance
  - ◆ Defaults may differ from statistics-focused alternatives
  - ◆ Very mature platform for machine learning
- Consider primary goals of analysis

# Regression with statsmodels

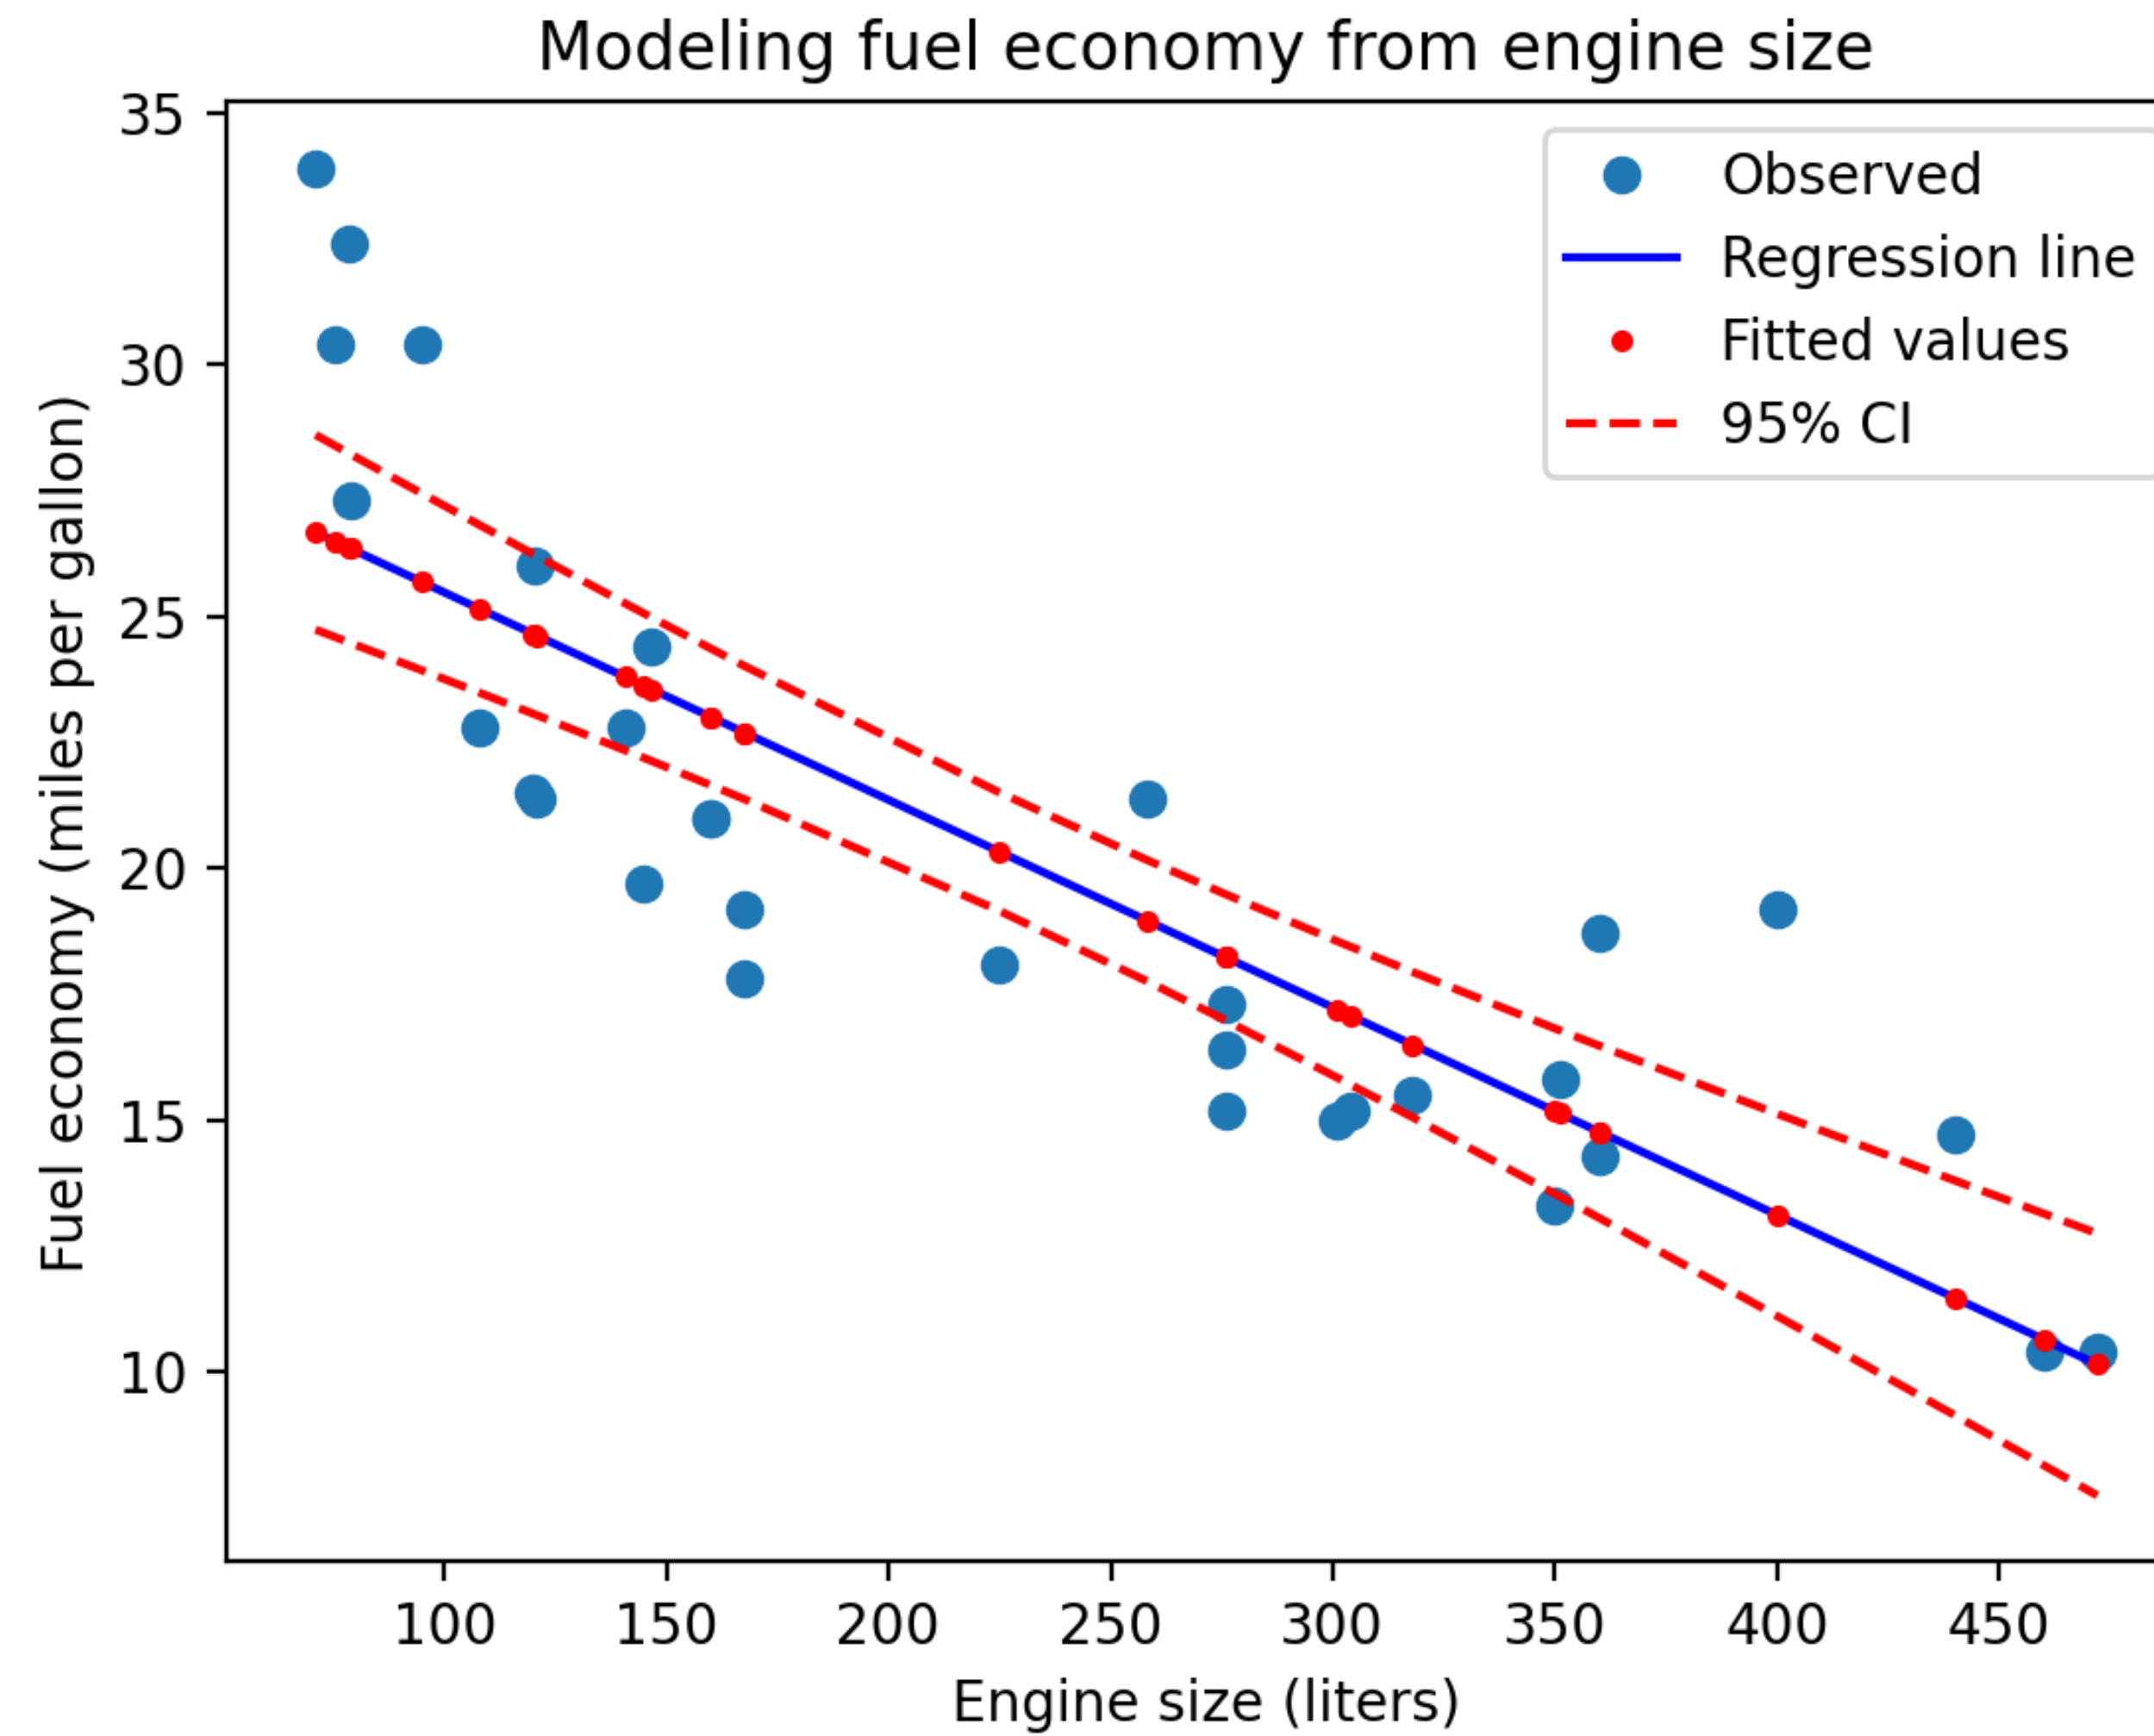
```

=====
                        OLS Regression Results
=====
Dep. Variable:          mpg      R-squared:                0.718
Model:                  OLS      Adj. R-squared:            0.709
Method:                 Least Squares      F-statistic:          76.51
Date:                  Tue, 13 Apr 2021      Prob (F-statistic):    9.38e-10
Time:                  17:21:57      Log-Likelihood:       -82.105
No. Observations:      32      AIC:                  168.2
Df Residuals:          30      BIC:                  171.1
                                Df Model:                1
                                Covariance Type:          nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      29.5999      1.230      24.070      0.000      27.088      32.111
disp      -0.0412      0.005      -8.747      0.000      -0.051      -0.032
=====
Omnibus:          3.368      Durbin-Watson:          0.986
Prob(Omnibus):    0.186      Jarque-Bera (JB):       3.049
Skew:             0.719      Prob(JB):               0.218
Kurtosis:         2.532      Cond. No.                558.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

# Regression with **statsmodels** (2)



# Classification with **sklearn**

```
In : from sklearn.model_selection import train_test_split
```

```
In : sonar = pd.read_csv("sonar.csv")
```

```
In : sonar_X = sonar.iloc[:,0:60]
```

```
In : sonar_y = sonar["Class"]
```

```
In : part = train_test_split(sonar_X, sonar_y, test_size=0.2)
```

```
In : X_train, X_test, y_train, y_test = part
```

```
In : model2.score(X_train, y_train) # accuracy (training)
```

```
Out: 0.8975903614457831
```

```
In : model2.score(X_test, y_test) # accuracy (testing)
```

```
Out: 0.8333333333333334
```

# Classification with **sklearn** (2)

```
In : from sklearn.metrics import classification_report
```

```
In : classification_report(y_test, model2.predict(X_test))
```

Out:

	precision	recall	f1-score	support
Mine	0.80	0.91	0.85	22
Rock	0.88	0.75	0.81	20
accuracy			0.83	42
macro avg	0.84	0.83	0.83	42
weighted avg	0.84	0.83	0.83	42



# Modeling considerations

- “All models are wrong, but some are useful.”
  - ◆ Consider what is the goal of the analysis
  - ◆ Evaluate a model by how *useful* it is for the goal
- Don't use a model you don't understand
  - ◆ Understand modeling assumptions
  - ◆ What kinds of models are appropriate?
- Follow best practices for model evaluation

# REVIEW