



Writing faster and better Matlab code

Preallocate arrays

- Create storage for array results before filling it.

```
% Matlab script -- compute sin(x) for values from 0:100000
```

```
clear all
```

```
% No preallocation
```

```
x = linspace(0, 100, 100000);
```

```
tic;
```

```
for i = 1:length(x)
```

```
    y(i) = sin(x(i));
```

```
end
```

```
t = toc;
```

```
fprintf('No preallocation.  Computed %d values in %f seconds\n', length(x), t)
```

```
clear all
```

```
% With preallocation
```

```
x = linspace(0, 100, 100000);
```

```
y = zeros(size(x));
```

```
tic;
```

```
for i = 1:length(x)
```

```
    y(i) = sin(x(i));
```

```
end
```

```
t = toc;
```

```
fprintf('With preallocation.  Computed %d values in %f seconds\n', length(x), t)
```

Preallocate arrays

```
>> PreallocateArrays
```

```
No preallocation.  Computed 100000 values in 0.045132 seconds
```

```
With preallocation.  Computed 100000 values in 0.008846 seconds
```

- Preallocation is about 5 -- 7x faster.
- With no preallocation, Matlab must continually spend time finding new storage for the growing output array.
- With preallocation, no time is wasted resizing the output array at each loop iteration.

Preallocate arrays of specific type

% Matlab script -- Draw white circle on black background

clear all

% No preallocation

Rows = 250; % Number of rows in image matrix.

Rc = Rows/2; % Where to place the image center (rows).

Cols = 250; % Number of cols in image matrix.

Cc = Cols/2; % Where to place the image center (cols).

Planes = 3; % Three color planes

Rad = 50; % Radius of circle to draw

tic;

for r = 1:Rows

for c = 1:Cols

for p = 1:Planes

if norm([r - Rc, c - Cc]) < Rad

% white circle is 255 on all planes

myimage(r, c, p) = uint8(255);

else

myimage(r, c, p) = uint8(0);

end

end

end

end

t = toc;

figure

image(myimage)

fprintf('No preallocation. Created matrix in %f seconds\n', t)

Notice uint8 cast
-- bad


Preallocate arrays of specific type

```
% Draw white circle on black background  
clear all
```

```
% With preallocation
```

```
Rows = 250;    % Number of rows in image matrix.  
Rc = Rows/2;   % Where to place the image center (rows).  
Cols = 250;    % Number of cols in image matrix.  
Cc = Cols/2;   % Where to place the image center (cols).  
Planes = 3;    % Three color planes  
Rad = 50;      % Radius of circle to draw  
myimage = zeros(Rows, Cols, Planes, 'uint8');  
tic;  
for r = 1:Rows  
    for c = 1:Cols  
        for p = 1:Planes  
            if norm([r - Rc, c - Cc]) < Rad  
                % Draw white circle  
                myimage(r, c, p) = 255;  
            end  
        end  
    end  
end  
t = toc;  
figure  
image(myimage)  
fprintf('With preallocation.  Created matrix in %f seconds\n', t)
```

Preallocate entire
uint8 array



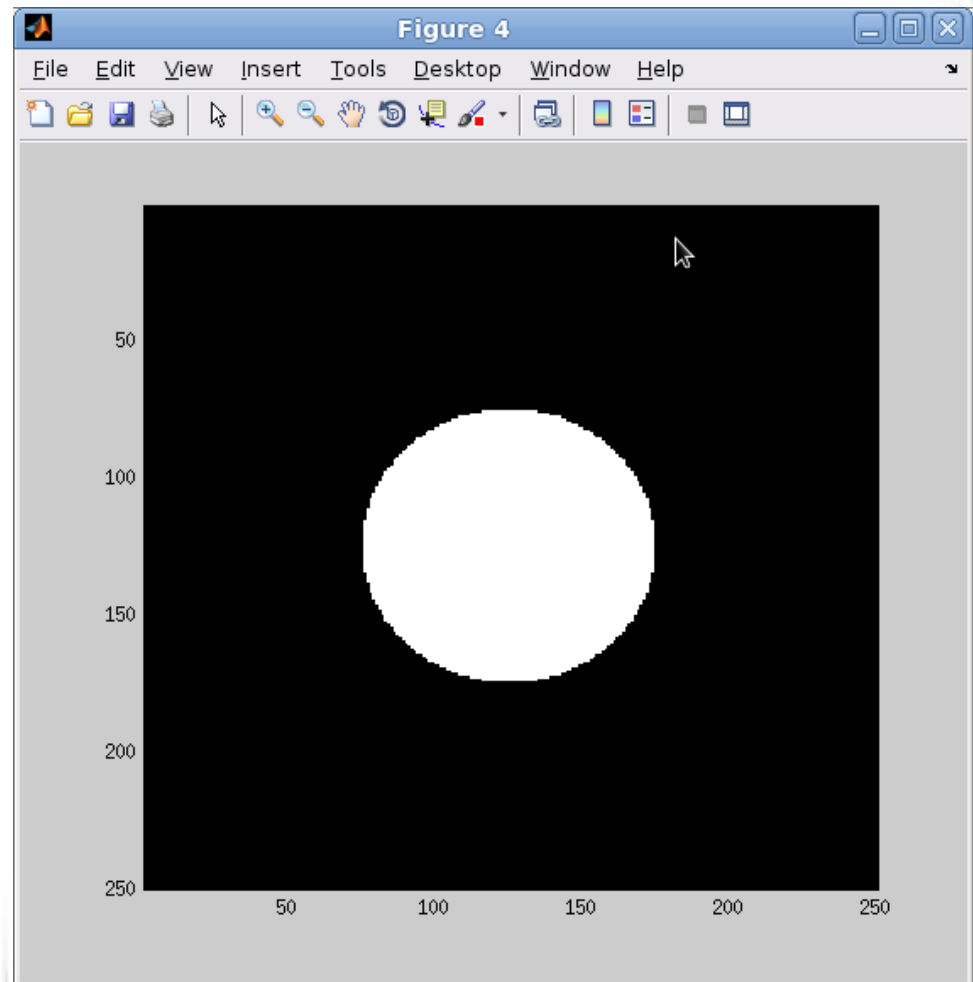
Comparison

```
>> PreallocateArrayType
```

No preallocation. Created matrix in 0.395536 seconds

With preallocation. Created matrix in 0.187897 seconds

- In this case, 2x speedup with preallocation.
- Matlab gives you the ability to create arrays of specific types using `zeros()`, `ones()`, `eye()`, etc.



Vectorize your code

- Vectorize = Avoid “for” loops when possible. Use built-in primitives which accept array inputs.

```
% Matlab script -- compute sin(x) for values from 0:100000
```

```
clear all
```

```
% No vectorization (but with preallocation)
```

```
x = linspace(0, 100, 100000);
```

```
y = zeros(size(x));
```

```
tic;
```

```
for i = 1:length(x)
```

```
    y(i) = sin(x(i));
```

```
end
```

```
t = toc;
```

```
fprintf('No vectorization.  Computed %d values in %f seconds\n', length(x), t)
```

```
clear all
```

```
% Vectorized
```

```
x = linspace(0, 100, 100000);
```

```
tic
```

```
y = sin(x);
```

```
t = toc;
```

```
fprintf('Vectorized.  Computed %d values in %f seconds\n', length(x), t)
```

Effects of vectorization

- Factor of 4 speedup with vectorized code.

```
>> Vectorize
```

```
No vectorization. Computed 100000 values in 0.009425 seconds
```

```
Vectorized. Computed 100000 values in 0.002328 seconds
```

- Effect is much more pronounced in Octave – 200x speedup

```
octave:3> Vectorize
```

```
No vectorization. Computed 100000 values in 1.140878 seconds
```

```
Vectorized. Computed 100000 values in 0.005408 seconds
```

- This is because Matlab has a very smart JIT – just in time compiler – which can recognize and vectorize simple for loops.

Use “meshgrid” to create {x, y} grids

```
>> [X, Y] = meshgrid(-2:2, -3:3)
```

X =

-2	-1	0	1	2
-2	-1	0	1	2
-2	-1	0	1	2
-2	-1	0	1	2
-2	-1	0	1	2
-2	-1	0	1	2
-2	-1	0	1	2


Y =

-3	-3	-3	-3	-3
-2	-2	-2	-2	-2
-1	-1	-1	-1	-1
0	0	0	0	0
1	1	1	1	1
2	2	2	2	2
3	3	3	3	3

- Meshgrid takes vectors, returns matrices.
- X, Y pairs form coordinate set $[x(i, j), y(i, j)]$
- Example:
 $[x(1,1), y(1,1)] = [-2, -3]$
- Example:
 $[x(5,5), y(5,5)] = [2, 1]$

Example of meshgrid use

- Create plot of function
$$z(x, y) = x e^{-(x^2 + y^2)}$$
- Non-vectorized algorithm would compute $z(i, j)$ using two for loops.
- Note: Time vs. memory trade-off.



```
u = linspace(-3, 3, 50);  
v = linspace(-3, 3, 50);  
[x, y] = meshgrid(u, v);  
z = x .* exp(-x.^2 - y.^2);  
surf(x, y, z)
```


Use logical indexing

- Logical indexing instead of for loops is a different aspect of vectorization.
- Here, we find all negative elements of A.
- Return is a mask.

```
>> A = randn(3,4)
```

```
A =
```

0.1832	0.3071	0.2614	-0.1461
-1.0298	0.1352	-0.9415	-0.5320
0.9492	0.5152	-0.1623	1.6821

```
>> lt = (A < 0)
```

```
lt =
```

0	0	0	1
1	0	1	1
0	0	1	0

```
>> A(lt)
```

```
ans =
```

-1.0298
-0.9415
-0.1623
-0.1461
-0.5320

Vectorized circle drawing program

% Matlab script -- Draw white circle on black background

clear all

% With preallocation, no vectorization.

Rows = 250; % Number of rows in image matrix.

Rc = Rows/2; % Where to place the image center (rows).

Cols = 250; % Number of cols in image matrix.

Cc = Cols/2; % Where to place the image center (cols).

Planes = 3; % Three color planes

Rad = 50; % Radius of circle to draw

myimage = zeros(Rows, Cols, Planes, 'uint8');

tic;

for r = 1:Rows

for c = 1:Cols

for p = 1:Planes

if norm([r - Rc, c - Cc]) < Rad

% Draw white circle

myimage(r, c, p) = 255;

end

end

end

end


t = toc;

figure

image(myimage)

fprintf('No vectorization, with preallocation. Created matrix in %f seconds\n', t)

Loop over all points, color point white if it lies inside the circle of radius Rad.



Vectorized circle drawing program

```
clear all
```

```
% With vectorization and preallocation.
```

```
Rows = 250; % Number of rows in image matrix.
```

```
Rc = Rows/2; % Where to place the image center (rows).
```

```
Cols = 250; % Number of cols in image matrix.
```

```
Cc = Cols/2; % Where to place the image center (cols).
```

```
Planes = 3; % Three color planes
```

```
Rad = 50; % Radius of circle to draw
```

```
myimage = zeros(Rows, Cols, Planes, 'uint8'); % Create black background
```

```
tic;
```

```
% Create row and col matrices
```

```
[r, c] = meshgrid(1:Rows, 1:Cols);
```

```
% find all [r, c] inside circle and create idx matrix
```

```
idx = (r - Rc).*(r - Rc) + (c - Cc).*(c - Cc) < Rad*Rad;
```

```
% Now draw circle a new 2D matrix using idx
```

```
mat = zeros(Rows, Cols);
```

```
mat(idx) = 255;
```

```
% Now copy this matrix into the 3D array myimage
```

```
for p = 1:Planes
```

```
    myimage(:, :, p) = mat;
```

```
end
```

```
t = toc;
```

```
figure
```

```
image(myimage)
```

```
fprintf('Vectorization and preallocation. Created matrix in %f seconds\n', t)
```

Use meshgrid to create plane of points. Then use logical indexing to select the points inside the circle.

Now color indexed points white.

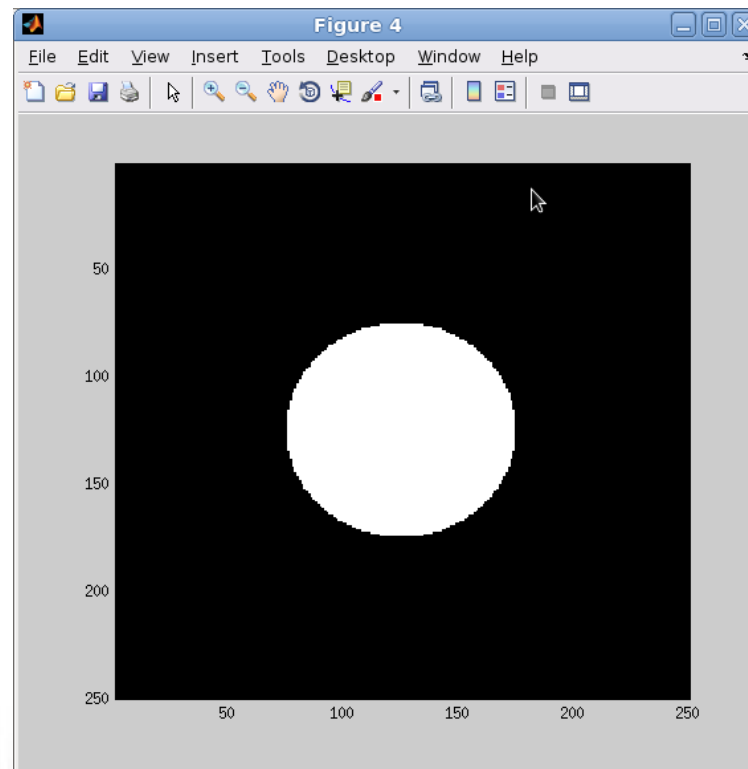
Performance comparison

```
>> PreallocateArraysTypeVectorized
```

No vectorization, with preallocation. Created matrix in 0.191326 seconds

Vectorization and preallocation. Created matrix in 0.008339 seconds

- > 20x improvement using vectorized code – in Matlab.



Use “find”

```
>> A = randn(7)
```

```
A =
```

-0.2725	-1.5771	-0.2991	-1.1564	-0.7982	-0.2938	-0.8655
1.0984	0.5080	0.0229	-0.5336	1.0187	-0.8479	-0.1765
-0.2779	0.2820	-0.2620	-2.0026	-0.1332	-1.1201	0.7914
0.7015	0.0335	-1.7502	0.9642	-0.7145	2.5260	-1.3320
-2.0518	-1.3337	-0.2857	0.5201	1.3514	1.6555	-2.3299
-0.3538	1.1275	-0.8314	-0.0200	-0.2248	0.3075	-1.4491
-0.8236	0.3502	-0.9792	-0.0348	-0.5890	-1.2571	0.3335

```
>> idx = find(abs(A)<0.1)
```

```
idx =
```

```
11  
16  
27  
28
```

“find” returns a
vector of indices



```
>> A(idx)
```

```
ans =
```

```
0.0335  
0.0229  
-0.0200  
-0.0348
```

Another find example

A =

-0.0942	-0.2883	1.0360
0.3362	0.3501	2.4245
-0.9047	-1.8359	0.9594

```
>> B = sqrt(A)
```

B =

0.0000 + 0.3070i	0.0000 + 0.5369i	1.0178 + 0.0000i
0.5798 + 0.0000i	0.5917 + 0.0000i	1.5571 + 0.0000i
0.0000 + 0.9511i	0.0000 + 1.3549i	0.9795 + 0.0000i

```
>> idx = find(imag(B) ~= 0)
```

idx =

1
3
4
6

```
>> A(idx)
```

ans =

-0.0942
-0.9047
-0.2883
-1.8359

Exploit functions which return indices

- Many functions return a value and an index.
- Min, max, intersect, union, sort, etc.

```
>> A = randn(1, 100);  
>> [x, i] = min(A)
```

```
x =  
  
-2.3299
```

```
i =  
  
72
```

```
>> A(i)
```

```
ans =  
  
-2.3299
```


Index using a vector of indices

```
>> y = randperm(20)
```

$$y =$$

10	18	12	13	8	11	16	20	6	19	7	3	15	9
5	4	14	1	2	17								

```
>> [ys, idx] = sort(y)
```

$$y_S =$$

1	2	3	4	5	6	7	8	9	10	11	12	13	14
15	16	17	18	19	20								

idx =

18	19	12	16	15	9	11	5	14	1	6	3	4	17
13	7	20	2	10	8								

```
>> y(idx)
```

ans =

1	2	3	4	5	6	7	8	9	10	11	12	13	14
15	16	17	18	19	20								

Use dedicated matrix constructors

- repmat

```
>> A = [0 0 1; 0 1 0; 1 0 0]
```

```
A =
```

0	0	1
0	1	0
1	0	0

```
>> B = repmat(A, 2, 2)
```

```
B =
```

0	0	1	0	0	1
0	1	0	0	1	0
1	0	0	1	0	0
0	0	1	0	0	1
0	1	0	0	1	0
1	0	0	1	0	0

Do the same thing using kron

- kron

```
octave:4> A = [0 0 1; 0 1 0; 1 0 0]  
A =
```

```
0  0  1  
0  1  0  
1  0  0
```

```
octave:5> R = ones(2,2)  
R =
```

```
1  1  
1  1
```

```
octave:6> kron(R, A)  
ans =
```

```
0  0  1  0  0  1  
0  1  0  0  1  0  
1  0  0  1  0  0  
0  0  1  0  0  1  
0  1  0  0  1  0  
1  0  0  1  0  0
```


Profile your code with tic/toc

```
>> profile on  
>> time_matmul
```

```
----- mymatmul -----  
N =      3, avg multiplication time =      0.00054825 sec  
N =      5, avg multiplication time =      0.00043700 sec  
N =     10, avg multiplication time =      0.00324100 sec  
N =     20, avg multiplication time =      0.02420525 sec  
N =     30, avg multiplication time =      0.07175375 sec  
N =     50, avg multiplication time =      0.30089525 sec  
N =    100, avg multiplication time =      2.31823850 sec  
N =    200, avg multiplication time =     21.46755650 sec
```

```
----- BLAS -----  
N =      3, avg multiplication time =      0.01171150 sec  
N =      5, avg multiplication time =      0.00001125 sec  
N =     10, avg multiplication time =      0.00205325 sec  
N =     20, avg multiplication time =      0.00002000 sec  
N =     30, avg multiplication time =      0.00002800 sec  
N =     50, avg multiplication time =      0.00008950 sec  
N =    100, avg multiplication time =      0.00025350 sec  
N =    200, avg multiplication time =      0.00083700 sec
```


```
My multiplication is 0(3.004)  
Matlabs multiplication is 0(1.751)
```

```
ans =
```

```
      3      5     10     20     30     50    100    200
```

```
>> profile report
```

The Mandelbrot Set




- Famous fractal.
- Fractal = Set with non-integer dimension.

Mandelbrot set mathematics

- Use iteration: $z_{n+1} = z_n^2 + c$ $z_0 = 0$
- z and c are complex numbers
- This iteration will either diverge to infinity, or stay bounded as c is varied.
- Make plot in complex plane of c :
 - If iteration remains bounded, plot point dark red
 - If iteration diverges, plot point blue.
- Note that z diverges to infinity when $|z| \geq 2$

$$z_{n+1} = z_n^2 + c$$



z walks to
infinity as
equation is
iterated

$\text{Re}(c)$

Iteration
remains
bounded

$\text{Im}(c)$

Algorithm

1. For $cr = -2:.01:2$
2. For $ci = -2:.01:2$
3. $z = 0$
4. For $cnt = 1:\text{max iterations}$
5. $z = z^2 + (cr + i*ci)$
6. If $\text{abs}(z) > 2$, iteration has diverged. Set $\text{pixel}(cr, ci)$ blue, then break to next ci iteration.
7. End cnt loop
8. Color $\text{pixel}(cr, ci)$ red
9. End ci loop
10. End cr loop

Example: Mandelbrot set

Non-vectorized

```
function Mandelbrot()  
    % This computes and plots the Mandelbrot set using a non-vectorized  
    % algorithm.  
  
    t = tic();      % Start performance timer  
  
    maxIterations = 255; % When to decide point will not escape to inf.  
    gridSize = 1000;    % This is length of each side of image matrix.  
  
    % Domain over which to compute Mandelbrot set.  
    xmin = -2;  
    xmax = 1;  
    ymin = -1;  
    ymax = 1;  
  
    % Start to setup z0 = x+iy values.  
    x = linspace( xmin, xmax, gridSize );  
    y = linspace( ymin, ymax, gridSize );  
  
    % Preallocate matrix to hold result of iteration.  
    % Note -- matrix rows are y, cols are x.  
    c = zeros(length(y), length(x));
```


Iterate over points
in complex plane

Iterate points to
check divergence

This version produces
more colors than
simple red/blue.







```
% Loop over all  $z = x+iy$  in complex plane
for xidx = 1:length(x)
    for yidx = 1:length(y)
        % Initialize z which will be iterated.
        z0 = complex(x(xidx), y(yidx));
        z = z0;
        for n = 0:maxIterations
            z = z*z + z0;
            if abs(z) <= 2
                % Still inside circle of radius 2
                c(yidx, xidx) = c(yidx, xidx) + 1;
            else
                % We've broken outside the circle -- stop iterating.
                break
            end
        end
        c(yidx, xidx) = log(c(yidx, xidx)+1); % This enhances the colors plotted
    end
end
cpuTime = toc( t ); % Stop performance timer

% Show plot
imagesc( x, y, c );
fprintf('Elapsed computation time = %f sec\n', cpuTime)


end
```

- Note three nested for loops.

Zoom sequence




Zoom into Mandelbrot set



- Play with colors and iteration count...

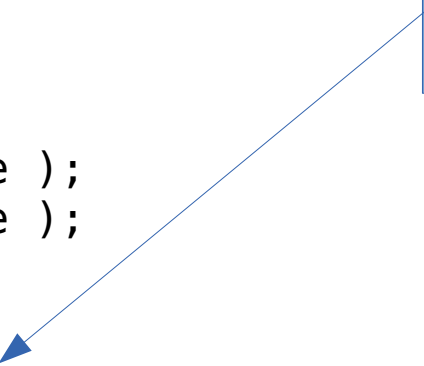
Deep zoom into Mandelbrot set



Vectorized Mandelbrot code

```
function Mandelbrot_vectorized()  
    % This computes and plots the Mandelbrot set using a fully-vectorized  
    % algorithm.  
  
    t = tic();      % Start performance timer  
  
    maxIterations = 50; % When to decide point will not escape to inf.  
    gridSize = 1000;    % This is length of each side of image matrix.  
  
    % Domain over which to compute Mandelbrot set.  
    xmin = -2;  
    xmax = 1;  
    ymin = -1;  
    ymax = 1;  
  
    % Setup z0 = x+iy values.  
    x = linspace( xmin, xmax, gridSize );  
    y = linspace( ymin, ymax, gridSize );  
  
    % Create initial z values  
    [xGrid,yGrid] = meshgrid( x, y );  
    z0 = complex(xGrid, yGrid);  
  
    % Preallocate matrix to hold result of iteration.  
    c = zeros( size(z0) );
```

Set up playing field
using meshgrid



This statement updates
entire complex plane at
once

```
% Calculate iteration
z = z0;
for n = 0:maxIterations
    z = z.*z + z0;
    % If the point is still inside the circle, add 1 to count.
    inside = (abs(z) <= 2); % Find indicies of points inside circle
    c = c + inside;
end
c = log(c+1); % Enhance colors plotted

cpuTime = toc( t ); % Stop timer

% Show plot
imagesc( x, y, c );
fprintf('Elapsed computation time = %f sec\n', cpuTime)


end
```

Use logical indexing
to find points
inside circle

Performance comparison

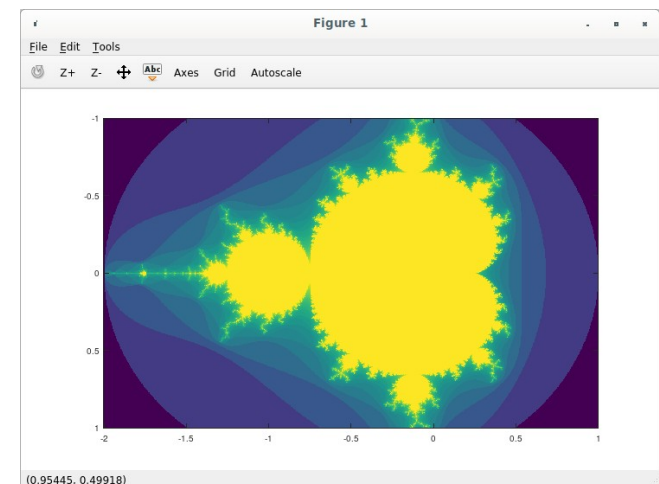
- Factor of 2 to 3 improvement for vectorized version (Matlab).

```
>> Mandelbrot  
Elapsed computation time = 11.827231 sec  
>> Mandelbrot_vectorized  
Elapsed computation time = 4.307583 sec
```



- Factor of 150 improvement for vectorized version (Octave).

```
octave:9> Mandelbrot  
Elapsed computation time = 253.655497 sec  
octave:10> Mandelbrot_vectorized  
Elapsed computation time = 1.591694 sec
```



Session summary

- Writing performant Matlab code
 - Preallocate arrays
 - Vectorize your code
 - Use specialized constructors for vectorization
 - Exploit functions which return index vectors
 - Use specialized matrix constructors like repmat, kron
- Mandelbrot set
- Mandelbrot set vectorized