

Math 7243 Machine Learning - Fall 2021

Instructor: He Wang

Test 1.

Student Name: _____/50

Rules and Instructions for Exams:

1. Unless otherwise specified, to receive full credits you must show **all** necessary work. The grading is based on your work shown. Only a final result from computer will receive zero point.
2. You need to finish the exam yourself. Any discussions with the other people will be considered as **academic dishonesty**. **Cheating, Unauthorized Collaboration, and Facilitating Academic Dishonesty are not allowed.** You can read a description of each here <http://www.northeastern.edu/osccr/academic-integrity-policy/>
3. This is an open exam. You are allowed to look at textbooks, and use a computer.
4. You are **not** allowed to discuss with any other people.
5. You are **not** allowed to ask questions on any internet platform.
6. For programming questions, if there is no specific instruction, you can only use numpy library. You should **not** use any build in function from Scikit-learn or StatsModels libraries.

1. (10 points) Calculate the **gradient** and **Hessian matrix** of the following functions and find the **argmin** _{θ} of each function. Here the norm $\|\cdot\|$ is the standard l_2 -norm. You can use any results in the lecture notes.
- (1) Let $\vec{b} \in \mathbb{R}^d$ and let $J(\vec{\theta}) = \|\vec{\theta} - \vec{b}\|^2$.

$$J(\vec{\theta}) = \|\vec{\theta} - \vec{b}\|^2 = (\vec{\theta} - \vec{b})^T (\vec{\theta} - \vec{b}) = \vec{\theta}^T \vec{\theta} - 2\vec{b}^T \vec{\theta} + \vec{b}^T \vec{b}$$

So, the gradient of $J(\vec{\theta})$ is

$$\Delta J(\vec{\theta}) = 2\vec{\theta} - 2\vec{b}$$

The Hessian matrix of $J(\vec{\theta})$ is

$$H(J(\vec{\theta})) = 2I_d$$

argmin _{θ} (J) is \vec{b}

- (2) Let $X \in \mathbb{R}^{n \times d}$ and $\vec{b} \in \mathbb{R}^d$. Suppose $\text{rank}(X) = d$. Let $F(\vec{\theta}) = \|X\vec{\theta}\|^2 + \vec{\theta}^T \vec{b}$.

$$J(\vec{\theta}) = \vec{\theta}^T X^T X \vec{\theta} + \vec{b}^T \vec{\theta}$$

So, the gradient of $J(\vec{\theta})$ is

$$\Delta J(\vec{\theta}) = 2X^T X \vec{\theta} + \vec{b}$$

The Hessian matrix of $J(\vec{\theta})$ is

$$H(J(\vec{\theta})) = 2X^T X$$

argmin _{θ} (J) is $-(2X^T X)^{-1}\vec{b}$

2. (10 points) In this question, you may use Python (with only numpy library) to solve the matrix equation. Consider the following data points

x_1	x_2	y
1.1	2	2.3
2.2	4	4.3
3.1	6	6.3
4.2	8	7.8
5.3	10	9.8

a). Fit a linear model $y = \theta_0 + \theta_1 x_1 + \theta_2 x_2$ to this dataset when the loss is $RSS = \|X\vec{\theta} - \vec{y}\|^2$. You should report the best fit function and the RSS cost value.

(1) We will use the formula $\beta = (X^T X)^{-1} X^T \vec{y}$ to do the calculation.
 The coefficient we get is $\vec{\theta} = \text{array}([0.6, -0.83333333, 1.35833333])$
 The RSS cost is $RSS(\theta) = \|X\vec{\theta} - \vec{y}\|^2 = 0.06667$

b). Fit a linear function to this dataset when the loss is the Ridge Loss $J(\theta) = \|X\vec{\theta} - \vec{y}\|^2 + \lambda(\theta_1^2 + \theta_2^2)$ with $\lambda = 1$ and with $\lambda = 10$. You should report the best fit function and the **RSS** cost value. (Warning: Do not put penalty on θ_0)

(2) We will use the formula $\beta = (X^T X + \lambda I)^{-1} X^T \vec{y}$ to do the calculation. However, we don't want to put penalty on β_0 .

We need to centralize our data by calculate $x' := x^{(i)} - \bar{x}$ and $y' := y^i - \bar{y}$

```

x.mean(axis=0) = array([3.18, 6. ])
y.mean()=6.1
x-x.mean(axis=0) = array([[ -2.08, -4. ], [ -0.98, -2. ], [ -0.08, 0. ], [ 1.02, 2. ], [ 2.12, 4. ]])
y-y.mean()=array([-3.8, -1.8, 0.2, 1.7, 3.7])

```

In this new centralized data, we have the linear model $y' = \beta_1 x'_1 + \beta_2 x'_2$ with no intersection.

So, now, we can use the formula to the new data (X, y) , and calculate the coefficient $\beta = (X^T X + \lambda I)^{-1} X^T \vec{y}$

$\bar{y} = 6.1$ and

So $y - \bar{y} = \beta_1(x_1 - \bar{x}_1) + \beta_2(x_2 - \bar{x}_2)$.

When $\lambda = 1$, $[\beta_1, \beta_2] = \text{array}([0.35998318, 0.71981341])$

The Cost value is 0.0967

When $\lambda = 10$ $[\beta_1, \beta_2] = \text{array}([0.31523096, 0.60886392])$

The Cost value is 1.009285

3. (10 points) Consider the data

$x^{(i)}$	0	0.2	0.4	0.6	0.8	1	1.2	1.4
$y^{(i)}$	5.1	6.4	6.1	8.2	9.5	8.6	12	14.8

The data file $\{\vec{x}^{(i)}, y^{(i)}\}$ for $i = 1, 2, \dots, n = 8$ is drawn (with noise) from

$$f(x) = \theta_0 + \theta_1 e^x$$

(1) Find a **closed formula** for parameters $\vec{\theta}$ to minimize the RSS loss

$$J(\vec{\theta}) = \sum_{i=1}^n (y^{(i)} - f(x^{(i)}))^2$$

The least squares solution is

$$\vec{\theta} = (X^T X)^{-1} X^T \vec{y}$$

where $X = \begin{bmatrix} 1 & \exp(\vec{x}^{(1)}) \\ 1 & \exp(\vec{x}^{(2)}) \\ \vdots & \vdots \\ 1 & \exp(\vec{x}^{(n)}) \end{bmatrix}$ and $\vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(n)} \end{bmatrix}$

(2) **Find the function** $f(x)$ fitting the data using the result in (1).

Define a new column $x = \text{np.exp}(x)$.

Then, use the formula $(X^T X)^{-1} X^T \vec{y}$ to find the $\vec{\theta} = [2.26789, 2.94362]$

$$f(x) = 2.26789 + 2.94362 e^x$$

The cost is 4.54324

4. (10 points) Consider the data

$x^{(i)}$	0	0.2	0.4	0.6	0.8	1	1.2	1.4
$y^{(i)}$	3.2	4.2	5.6	5.2	7.7	8.8	13.9	18.7

The data file $\{\vec{x}^{(i)}, y^{(i)}\}$ for $i = 1, 2, \dots, n = 8$ is drawn (with noise) from the function:

$$g(x) = \theta_0 + e^{\theta_1 x}.$$

Fit the data to the function $g(x)$ by minimizing the RSS loss

$$J(\vec{\theta}) = \sum_{i=1}^n (y^{(i)} - g(x^{(i)}))^2.$$

(1) Find the **gradient** of the cost function $J(\vec{\theta})$.

$$J(\vec{\theta}) = \sum_{i=1}^n (y^{(i)} - g(x^{(i)}))^2 = \sum_{i=1}^n (y^{(i)} - \theta_0 - e^{\theta_1 x})^2.$$

$$\frac{\partial}{\partial \theta_0} = -2 \sum_{i=1}^n (y^{(i)} - g(x^{(i)}))$$

$$\frac{\partial}{\partial \theta_1} = -2 \sum_{i=1}^n (y^{(i)} - g(x^{(i)})) \theta_1 e^{\theta_1 x}$$

(2) Write down the update formula for gradient decent using α for the learning rate.

$$\theta_0^{next} = \theta_0 + \alpha 2 \sum_{i=1}^n (y^{(i)} - g(x^{(i)}))$$

$$\theta_1^{next} = \theta_1 + 2 \sum_{i=1}^n (y^{(i)} - g(x^{(i)})) \theta_1 e^{\theta_1 x}$$

(3) Use gradient decent(GD) to find θ_* to minimize $J(\vec{\theta})$. You should try different learning rates and recording the cost function values to see what is the best α . Turn in any associated computations, your learning rate, cost values, and the parameters.

The model I used is $f(x) = 3 + e^{2x} + noise$

5. Consider the categorical learning problem consisting of a data set with two labels:

Label 1: (contains 6 points)

X_1	0.2	0.6	2	2.6	3.1	3.8
X_2	3.4	1.8	2	2.7	3.5	1.5

Label 2: (contains 5 points)

X_1	-0.7	-2.1	-2.5	-3	-3.9
X_2	-2.9	-2.8	-1.3	-2	-1.5

(1) (7 points) For each label above, the data follow a multivariate normal distribution $\text{Normal}(\mu_i, \Sigma)$ where the covariance Σ is the same for both labels. Fit a pair of LDA functions to the labels by computing the covariances Σ , means μ_i , and proportion ϕ of data. You may use Python (with only numpy library)

(a) You should report the values for ϕ , μ_i and Σ .

$$\begin{aligned}\mu_1 &= \begin{bmatrix} 2.05 \\ 2.48 \end{bmatrix} \\ \mu_2 &= \begin{bmatrix} -2.44 \\ -2.1 \end{bmatrix} \\ \Sigma &= \frac{1}{11-2} \sum_{i=1}^1 1(x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T = \begin{bmatrix} 1.732 & -0.4025 \\ -0.4025 & 0.636 \end{bmatrix} \\ \phi &= 6/11\end{aligned}$$

(b) Give the **formula for the line** forming the decision boundary.

The line forming the discretion boundary is $\vec{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$ such that $\log\left(\frac{P(\text{label} = 1|\vec{x})}{P(\text{label} = 2|\vec{x})}\right)$

$$P(\text{label} = k|\vec{x}) = \frac{P(\vec{x}|y = k)P(y = k)}{P(\vec{x})}$$

Simplify the calculation with constants, we have the equality

$$\vec{x}^T \Sigma^{-1} \mu_1 - \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \log(\phi_1) = \vec{x}^T \Sigma^{-1} \mu_2 - \frac{1}{2} \mu_2^T \Sigma^{-1} \mu_2 + \log(\phi_2)$$

Plug in the information from (1) we have the line

$$5.21x_1 + 10.68x_2 = 0.62$$

(2) (3 points) For each label above, use **logistic regression** to classify the data. You should report the **logistic function** $p(Y = 1|\vec{x}) = \frac{1}{1 + e^{-\theta^T \vec{x}}}$ and the **formula for the line** forming the decision boundary. (In this question, you can use any Python library including Scikit-learn.)

$$p(y = 1|\vec{x}) = \frac{1}{(1 + \exp(-(0.2237 + 0.6839x_1 + 0.8666x_2)))}$$

The decision boundary is $0.2237 + 0.6839x_1 + 0.8666x_2 = 0$

5. (continue)

(3) (2 bonus points) Find the probability $P(y = 1|\vec{x})$ for a test point $\vec{x} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ for both LDA model and the logistics model in the above two questions.

For LDA

$$P(\text{label} = 1|\vec{x}) = \frac{P(\vec{x}|y=1)P(y=1)}{P(\vec{x})} = \frac{P(\vec{x}|y=1)P(y=1)}{P(\vec{x}|y=1)P(y=1) + P(\vec{x}|y=2)P(y=2)}$$

$$= \frac{\phi_1 \frac{1}{2\pi|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\vec{x} - \mu_1)^T \Sigma^{-1}(\vec{x} - \mu_1)\right)}{\phi_1 \frac{1}{2\pi|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\vec{x} - \mu_1)^T \Sigma^{-1}(\vec{x} - \mu_1)\right) + \phi_2 \frac{1}{2\pi|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\vec{x} - \mu_2)^T \Sigma^{-1}(\vec{x} - \mu_2)\right)}$$

When $\vec{x} = (0, 0)$,

$$P(\text{label} = 1|\vec{x}) = 0.3486$$

For Logistic regression: $p(y = 1|\vec{x}) = \frac{1}{(1 + \exp(-(0.2237 + 0.6839x_1 + 0.8666x_2)))} = 0.55$ when $\vec{x} = (0, 0)$

(4) (2 bonus points) Find the boundary using the QDA method. (You may use a computer, but only with numpy library)

We assume the covariance Σ_1 and Σ_2 for each label are different. In this case,

$$\Sigma_1 = \frac{1}{6-1} \sum_{i=1}^6 (x^{(i)} - \mu_{y(i)})(x^{(i)} - \mu_{y(i)})^T = \begin{bmatrix} 9.995 & -1.215 \\ -1.215 & 3.883 \end{bmatrix}$$

$$\Sigma_2 = \frac{1}{5-1} \sum_{i=1}^5 (x^{(i)} - \mu_{y(i)})(x^{(i)} - \mu_{y(i)})^T = \begin{bmatrix} 5.592 & -2.61 \\ -2.61 & 2.14 \end{bmatrix}$$

Simplify the calculation with constants, we have the equality

$$-\frac{1}{2} \log |\Sigma_1| - \frac{1}{2}(\vec{x} - \mu_1)^T \Sigma_1^{-1}(\vec{x} - \mu_1) + \log(\phi_1) = -\frac{1}{2} \log |\Sigma_2| - \frac{1}{2}(\vec{x} - \mu_2)^T \Sigma_2^{-1}(\vec{x} - \mu_2) + \log(\phi_2)$$

Plug in the information from (1) we have the quadratic curve